


# Chapter 3: SQL





# Contents

- Data Definition
- Basic Query Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Complex Queries
- Views
- Modification of the Database
- Joined Relations\*\*





# History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
  - SQL-86, SQL-89, SQL-92, SQL:1999, SQL:2003





# History<sup>Cont.</sup>

- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
  - Not all examples here may work on your particular system



# Data Definition Language

- Allows the specification of information about each relation, including:
  - The schema for each relation.
  - The domain of values associated with each attribute.
  - Integrity constraints
  - The set of indices to be maintained for each relations.
  - Security and authorization information for each relation.
  - The physical storage structure of each relation on disk.



# Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).



# Domain Types in SQL<sup>Cont.</sup>

- **numeric(p,d).** Fixed point number, with user-specified precision of  $p$  digits, with  $d$  digits to the right of decimal point.
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(n).** Floating point number, with user-specified precision of at least  $n$  digits.
- More are covered in Chapter 4.



# Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,  
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

- *r* is the name of the relation
- each *A<sub>i</sub>* is an attribute name in the schema of relation *r*
- *D<sub>i</sub>* is the data type of values in the domain of attribute *A<sub>i</sub>*





# Create Table Construct **Cont.**

- Example:

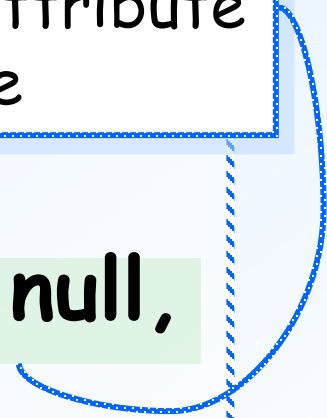
```
create table branch
```

```
(branch_name char(15) not null,
```

```
branch_city char(30),
```

```
assets integer)
```

null values is not  
allowed in attribute  
branch\_name



# Integrity Constraints

- **not null**
- **primary key** ( $A_1, \dots, A_n$ )

automatically  
ensures **not null** in  
SQL-92 onwards

Example: Declare *branch\_name* as the primary key for *branch*

```
create table branch
    (branch_name    char(15),
     branch_city    char(30),
     assets         integer,
     primary key (branch_name))
```




# Drop Table Constructs

- The **drop table** command deletes all information about the dropped relation from the database.

**drop table r**

Deletes not only all tuples of r, but also the schema for r

More drastic than



**delete table r**

Retains relation r, but deletes all tuples in r



# Alter Table Constructs

- The **alter table** command is used to add attributes to an existing relation:

```
alter table r add A D
```

- where *A* is the name of the attribute to be added to relation *r*; *D* is the domain of *A*.

# All tuples in the relation are assigned *null* as the value for the new attribute



# Alter Table Constructs Cont.

- The **alter table** command can also be used to drop attributes of a relation:

```
alter table r drop A
```

- where *A* is the name of an attribute of relation *r*

# Dropping of attributes → not supported by many databases

# Basic Query Structure

- SQL is based on relational algebra and relational calculus
- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

Attributes

Relations

Predicate



# Basic Query Structure Cont.

- equivalent to the relational algebra expression

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

- The result of an SQL query is a relation



# The *select* Clause

- The **select** clause list the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra

## Example

■ find the names of all branches in the *loan* relation

```
select branch_name  
from loan
```

corresponds to

```
 $\Pi_{branch\_name}(loan)$ 
```





# The select Clause **Cont.**

- SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)

## **Example**

*Branch\_Name*  $\equiv$  *BRANCH\_NAME*  $\equiv$  *branch\_name*

- SQL allows duplicates in relations as well as in query results





# The select Clause **Cont.**

- To force the elimination of duplicates, insert the keyword **distinct** after select.

## Example

■ Find the names of all branches in the *loan* relations, and remove duplicates

```
select distinct branch_name  
from loan
```






# The select Clause\*\*\*Cont.

- The keyword **all** specifies that duplicates not be removed.

```
select all branch_name  
from loan
```

- Default is **all**





# The select Clause **Cont.**


- An asterisk in the select clause denotes “all attributes”

```
select *  
from loan
```



Equal to

```
select loan_number, branch_name, amount  
from loan
```



# The *select* Clause Cont.

- The **select** clause can contain arithmetic expressions involving the operation,  $+$ ,  $-$ ,  $*$ , and  $/$ , and operating on constants or attributes of tuples

```
select loan_number, branch_name, amount*100  
from loan
```

- would return a relation that is the same as the *loan* relation, except that the value of the attribute *amount* is multiplied by 100



# The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra

## Example

■ find all loan number for loans made at the Perryridge branch with loan amounts greater than \$1200

```
select loan_number  
from loan  
where branch_name = 'Perryridge' and amount > 1200
```



# The where Clause **Cont.**

- SQL use the logical connectives **and**, **or**, **not** in the where clause
- Operands of logical connectives can be expressions involving the comparison operation  $<$ ,  $>$ ,  $<=$ ,  $>=$ ,  $=$ ,  $<>$
- Comparisons can be applied to results of arithmetic expressions, strings, special types such as date types



# The where Clause **Cont.**

- SQL includes a **between** comparison operator

## **Example**

- Find the loan number of those loans with loan amounts between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )

```
select loan_number  
from loan  
where amount between 90000 and 100000
```





# The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.

## Example

- Find the Cartesian product  
*borrower*  $\times$  *loan*

```
select *  
from borrower, loan
```



## Example

- Find the name, loan number and loan amount of all customers having a loan at the Perryridge branch

```
select customer_name, borrower.loan_number, amount
from borrower, loan
where borrower.loan_number = loan.loan_number and
        branch_name = 'Perryridge'
```

# The Rename Operation

- The SQL allows renaming attributes in the select clause using the **as** clause:

*old-name as new-name*

## Example

- Find the name, loan number and loan amount of all customers; rename the column name *loan\_number* as *loan\_id*

```
select customer_name,  
       borrower.loan_number as loan_id, amount  
from borrower, loan  
where borrower.loan_number = loan.loan_number
```



# Tuple Variables\_

- Tuple variables are defined in the **from** clause via the use of the **as** clause.

## Example

- Find the customer names and their loan numbers and loan amounts for all customers having a loan at some branch

```
select customer_name, T.loan_number, S.amount  
from borrower as T, loan as S  
where T.loan_number = S.loan_number
```



## Example

- Find the names of all branches that have greater assets than some branch located in Brooklyn

```
select distinct T.branch_name  
from branch as T, branch as S  
where T.assets > S.assets and S.branch_city = 'Brooklyn'
```

# Keyword **as** is optional and may be omitted  
*borrower as T*  $\equiv$  *borrower T*





# String Operations

- SQL includes a string-matching operator for comparisons on character strings.
- The operator “like” uses patterns that are described using two special characters:
  - percent (%). The % character matches any substring that have any length (can be 0).
  - underscore (\_). The \_ character matches any character.



## Example

- Find the names of all customers whose street includes the substring “Main”.

```
select customer_name  
from customer  
where customer_street like '% Main%'
```

```
# Match the string “80%”  
like '80\%' escape '\'
```



# String Operations Cont.

- SQL supports a variety of string operations such as
  - concatenation (using “||”)
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.



# Ordering the Display of Tuples

- The **order by** clause causes the tuples in the result of a query to appear in sorted order

## Example

- List in alphabetic order the names of all customers having a loan in Perryridge branch

```
select distinct customer_name
from    borrower, loan
where   borrower loan_number = loan.loan_number and
        branch_name = 'Perryridge'
order by customer_name
```

# Ordering the Display of Tuples Cont.

- We may specify **desc** for descending order or **asc** for ascending order
- Ordering can be performed on multiple attributes

```
select *  
from    loan  
order by amount desc, loan_number asc
```

# For each attribute, ascending order is the default





# Duplicates

- **Multiset** versions of some of the relational algebra operators – given multiset relations  $r_1$  and  $r_2$ :
  1.  $\sigma_{\theta}(r_1)$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$ , and  $t_1$  satisfies selections  $\sigma_{\theta}$ , then there are  $c_1$  copies of  $t_1$  in  $\sigma_{\theta}(r_1)$ .
  2.  $\Pi_A(r)$ : For each copy of tuple  $t_1$  in  $r_1$ , there is a copy of tuple  $\Pi_A(t_1)$  in  $\Pi_A(r_1)$  where  $\Pi_A(t_1)$  denotes the projection of the single tuple  $t_1$ .
  3.  $r_1 \times r_2$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$  and  $c_2$  copies of tuple  $t_2$  in  $r_2$ , there are  $c_1 \times c_2$  copies of the tuple  $t_1, t_2$  in  $r_1 \times r_2$

# Duplicates Cont.

- Example: Suppose multiset relations  $R_1 (A, B)$  and  $R_2 (C)$  are as follows:

$$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then  $\Pi_B(r_1)$  would be  $\{(a), (a)\}$ , while  $\Pi_B(r_1) \times r_2$  would be

$$\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$$

- SQL duplicate semantics:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

is equivalent to the *multiset* version of the expression

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$



# Set Operations

- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations  $\cup$ ,  $\cap$ ,  $-$ .
- Each of the above operations automatically eliminates duplicates



# Set Operations **Cont.**

- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  $r$  **union all**  $s$
  - $\min(m, n)$  times in  $r$  **intersect all**  $s$
  - $\max(0, m - n)$  times in  $r$  **except all**  $s$

## Example

- Find all customers who have a loan, an account, or both

```
(select customer_name from depositor)  
union  
(select customer_name from borrower)
```

Find all customers who have both a loan and an account

```
(select customer_name from depositor)  
intersect  
(select customer_name from borrower)
```

Find all customers who have an account but no loan.

```
(select customer_name from depositor)  
except  
(select customer_name from borrower)
```





# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values





## Example

- Find the average account balance at the Perryridge branch

```
select avg (balance)
from account
where branch_name = 'Perryridge'
```

- Find the number of tuples in the *customer* relation.

```
select count (*)
from customer
```

Count the number of  
tuples in a relation

- Find the number of depositors in the bank.

```
select count (distinct customer_name)
from depositor
```

eliminate duplicates


# Aggregate Functions – Group By

- Find the number of depositors for each branch

```
select branch_name, count (distinct customer_name)
from depositor, account
where depositor.account_number = account.account_number
group by branch_name
```

# Attributes in **select** clause outside of aggregate functions must appear in **group by** list





# Aggregate Functions – Having

- Find the names of all branches where the average account balance is more than \$1,200

```
select branch_name, avg (balance)
from account
group by branch_name
having avg (balance) > 1200
```

- # predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups



# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The predicate **is null** can be used to check for null values

## Example

- Find all loan number which appear in the loan relation with null values for amount

```
select loan_number  
from loan  
where amount is null
```



# Null Values Cont.


- The result of any arithmetic expression involving *null* is *null*

$5 + \text{null}$  returns null

- Any comparison with *null* returns *unknown*

$5 < \text{null}$  or  $\text{null} \neq \text{null}$  or  $\text{null} = \text{null}$

$\text{null is null}$   true

$\text{null} = \text{null}$   unknown



# Null Values **Cont.**

- Three-valued logic using the truth value *unknown*:
  - OR: (*unknown* **or** *true*) = *true*,  
(*unknown* **or** *false*) = *unknown*  
(*unknown* **or** *unknown*) = *unknown*
  - AND: (*true* **and** *unknown*) = *unknown*,  
(*false* **and** *unknown*) = *false*,  
(*unknown* **and** *unknown*) = *unknown*
  - NOT: (**not** *unknown*) = *unknown*



# Null Values<sup>Cont.</sup>


- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*
- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes

## Example

■ Total all loan amounts

```
select sum (amount )  
from loan
```

- # Above statement ignores null amounts
- # Result is *null* if there is no non-null amount



# Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries.
- A **subquery** is a **select-from-where** expression that is nested within another query.
- A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality



# Set Membership

Set is produced by subquery

- Find all customers who have both an account and a loan at the bank.

Test for set membership

```
select distinct customer_name
from borrower
where customer_name in (select customer_name
                        from depositor)
```

Find all customers who have a loan at the bank but do not have an account at the bank

Test for absence of set membership

```
select distinct customer_name
from borrower
where customer_name not in (select customer_name
                           from depositor)
```

# Example Query

- Find all customers who have both an account and a loan at the Perryridge branch

```
select distinct customer_name
from borrower, loan
where borrower.loan_number = loan.loan_number and
      branch_name = 'Perryridge' and
      (branch_name, customer_name) in
      (select branch_name, customer_name
       from depositor, account
       where depositor.account_number =
             account.account_number )
```

# Set Comparison

- Find all branches that have greater assets than some branch located in Brooklyn.

```
select distinct T.branch_name
from branch as T, branch as S
where T.assets > S.assets and
      S.branch_city = 'Brooklyn'
```

Same query using > **some** clause

```
select branch_name
from branch
where assets > some
      (select assets
       from branch
       where branch_city = 'Brooklyn')
```

# Definition of Some Clause

- **F <comp> some r**

–  $\exists t \in r$  such that  $(F \text{ <comp> } t)$ , Where  
<comp> can be: <, ≤, >, =, ≠

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$      $(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 5 \\ \hline \end{array}) = \text{true}$

**(= some) ≡ in**

However, **(≠ some) ≠ not in**

# Example Query

- Find the names of all branches that have greater assets than all branches located in Brooklyn

```
select branch_name
from branch
where assets > all
      (select assets
       from branch
       where branch_city = 'Brooklyn')
```



# Definition of all Clause

- $F <\text{comp}> \mathbf{all} \ r \Leftrightarrow \forall t \in r (F <\text{comp}> t)$

$$(5 < \mathbf{all} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 < \mathbf{all} \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 = \mathbf{all} \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \neq \mathbf{all} \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \mathbf{all}) \equiv \mathbf{not\ in}$

However,  $(= \mathbf{all}) \neq \mathbf{in}$



# Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$

# Example Query

- Find all customers who have both an account and a loan at the bank.

```
select distinct customer_name
from borrower
where exists (select *
              from depositor
              where depositor.customer_name =
                    borrower.customer_name )
```

```
select distinct customer_name
from borrower
where customer_name in (select customer_name
                        from depositor )
```





# Example Query

- Find all customers who have an account at all branches located in Brooklyn.

```
select distinct S.customer_name
from depositor as S
where not exists (
    (select branch_name
     from branch
     where branch_city = 'Brooklyn')
    except
    (select R.branch_name
     from depositor as T, account as R
     where T.account_number = R.account_number and
           S.customer_name = T.customer_name ))
```

# Test for Absence of Duplicate Tuples\*\*\*\*\*

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.

■ Find all customers who have at most one account at the Perryridge branch

```
select T.customer_name
from depositor as T
where unique (
    select R.customer_name
    from account, depositor as R
    where T.customer_name = R.customer_name and
           R.account_number = account.account_number and
           account.branch_name = 'Perryridge')
```

# Example Query

- Find all customers who have at least two accounts at the Perryridge branch.

```
select distinct T.customer_name
from depositor as T
where not unique (
  select R.customer_name
  from account, depositor as R
  where T.customer_name = R.customer_name and
    R.account_number = account.account_number and
    account.branch_name = 'Perryridge')
```

# Variable from outer level is known as a **correlation variable**



# Derived Relations

- SQL allows a subquery expression to be used in the **from** clause

■ Find the average account balance of those branches where the average account balance is greater than \$1200

```
select branch_name, avg_balance
from (select branch_name, avg (balance)
      from account
      group by branch_name )
      as branch_avg ( branch_name, avg_balance )
where avg_balance > 1200
```



# With Clause

- The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.

■ Find all accounts with the maximum balance

```
with max_balance (value) as  
    select max (balance)  
    from account
```

```
select account_number  
from account, max_balance  
where account.balance = max_balance.value
```


# Complex Queries using With Clause

- Find all branches where the total account deposit is greater than the average of the total account deposits at all branches.

```
with branch_total (branch_name, value) as  
    select branch_name, sum (balance)  
    from account  
    group by branch_name
```

```
with branch_total_avg (value) as  
    select avg (value)  
    from branch_total
```

```
select branch_name  
from branch_total, branch_total_avg  
where branch_total.value >= branch_total_avg.value
```





# Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know a customer's name, loan number and branch name, but has no need to see the loan amount. This person should see a relation described, in SQL.

```
(select customer_name, borrower.loan_number,  
                                             branch_name  
from borrower, loan  
where borrower.loan_number = loan.loan_number )
```



# Views **Cont.**

- A view is defined using the **create view** statement which has the form

```
create view v as < query expression >
```

- where <query expression> is any legal SQL expression. The view name is represented by v





Cont.

# Views

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- When a view is created, the query expression is stored in the database; the expression is substituted into queries using the view

# Example Query

- A view consisting of branches and their customers

```
create view all_customer as
(select branch_name, customer_name
 from depositor, account
 where depositor.account_number =
                                account.account_number )

union
(select branch_name, customer_name
 from borrower, loan
 where borrower.loan_number = loan.loan_number)
```

# Example Query


Find all customers of the Perryridge branch

```
select customer_name  
from all_customer  
where branch_name = 'Perryridge'
```



# Views Defined Using Other Views

- One view may be used in the expression defining another view
- A view relation  $v_1$  is said to *depend directly* on a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$



# Views Defined Using Other Views

- A view relation  $v_1$  is said to *depend on* view relation  $v_2$  if either  $v_1$  depends directly to  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
- A view relation  $v$  is said to be *recursive* if it depends on itself.





# View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view  $v_1$  be defined by an expression  $e_1$  that may itself contain uses of view relations.



# View Expansion **Cont.**

- View expansion of an expression repeats the following replacement step:

**repeat**

Find any view relation  $vi$  in  $e1$

Replace the view relation  $vi$  by the expression defining  $vi$

**until** no more view relations are present in  $e1$

- As long as the view definitions are not recursive, this loop will terminate



# Modification of the Database

## Deletion

- Delete all account tuples at the Perryridge branch

```
delete from account  
where branch_name = 'Perryridge'
```

- Delete all accounts at every branch located in the city 'Needham'.

```
delete from account  
where branch_name in  
    (select branch_name  
     from branch  
     where branch_city = 'Needham')
```



# Example Query

- Delete the record of all accounts with balances below the average at the bank.

```
delete from account
where balance < (select avg (balance )
                 from account )
```

- ⚠ Problem: as we delete tuples from deposit, the average balance changes
- ⚠ Solution used in SQL:
  1. First, compute **avg** balance and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

# Modification of the Database Insertion

- Add a new tuple to *account*

```
insert into account  
values ('A-9732', 'Perryridge', 1200)
```



Or Equivalently

```
insert into account (branch_name, balance, account_number)  
values ('Perryridge', 1200, 'A-9732')
```

- Add a new tuple to *account* with *balance* set to null

```
insert into account  
values ('A-777', 'Perryridge', null)
```

# Example Query

- Provide as a gift for all loan customers of the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account

**insert into account**

```
select loan_number, branch_name, 200
from loan
where branch_name = 'Perryridge'
```

**insert into depositor**

```
select customer_name, loan_number
from loan, borrower
where branch_name = 'Perryridge'
and loan.account_number = borrower.account_number
```

# Modification of the Database Updates

- Increase all accounts with balances over \$10,000 by 6%, all other accounts receive 5%.
  - Write two **update** statements:

1.

```
update account
set balance = balance * 1.06
where balance > 10000
```

2.

```
update account
set balance = balance * 1.05
where balance ≤ 10000
```


- The order is important
- Can be done better using the **case** statement (next slide)



# ✓ ✓ Case Statement for Conditional Updates

- Same query as before: Increase all accounts with balances over \$10,000 by 6%, all other accounts receive 5%.

```
update account
  set balance = case
    when balance <= 10000
      then balance * 1.05
    else  balance * 1.06
  end
```



# Update Through Views

- Create a view of all loan data in the *loan* relation, hiding the *amount* attribute

```
create view loan_branch as  
select loan_number, branch_name  
from loan
```

- Add a new tuple to *branch\_loan*

```
insert into loan_branch  
values ('L-37', 'Perryridge')
```

# represented by the insertion of the tuple  
( 'L-37', 'Perryridge', null ) into the *loan* relation

# Update Through Views **Cont.**

- Some updates through views are impossible to translate into updates on the database relations

```
create view v as
```

```
    select loan_number, branch_name, amount  
    from loan  
    where branch_name = 'Perryridge'
```

```
insert into v values ( 'L-99', 'Downtown', 23)
```



# Update Through Views

- Others cannot be translated uniquely

```
insert into all_customer  
values ('Perryridge', 'John')
```

- Have to choose loan or account, and create a new loan/account number!
- Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation







# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join



# Joined Relations Cont.

- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated

<i>Join types</i>
inner join
left outer join
right outer join
full outer join

<i>Join Conditions</i>
natural
on <predicate>
using ( $A_1, A_1, \dots, A_n$ )

# Joined Relations Datasets for Examples

■ Relation *loan*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

*loan*

Relation *borrower*

<i>customer_name</i>	<i>loan_number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

*borrower*

# borrower information missing for L-260 and loan information missing for L-155



# Joined Relations – Examples

*loan inner join borrower on  
loan.loan\_number = borrower.loan\_number*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>	<i>loan_number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230

*loan left outer join borrower on  
loan.loan\_number = borrower.loan\_number*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>	<i>loan_number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	<i>null</i>	<i>null</i>




# Joined Relations – Examples

*loan natural inner join borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

*loan natural right outer join borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	<i>null</i>	<i>null</i>	Hayes



# Joined Relations – Examples

*loan full outer join borrower using (loan\_number)*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	<i>null</i>	<i>null</i>	Hayes

Find all customers who have either an account or a loan (but not both) at the bank.

```
select customer_name  
from (depositor natural full outer join borrower)  
where account_number is null or loan_number is null
```



# Conclusions






# Questions?







# End of Chapter 3

